



DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

Module 3



DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

Deploying Your Prototype

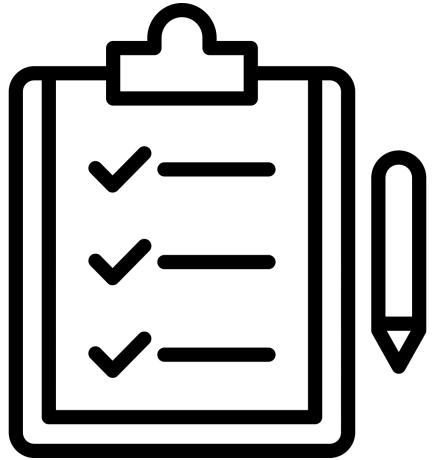


DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

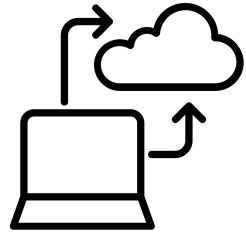
Choosing the Right Deployment Strategy

MVP Build Plan



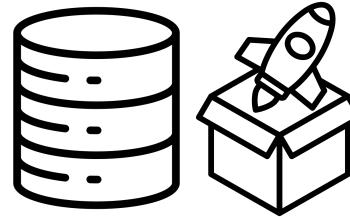
1. Getting data into Snowflake
2. Cleanup Data
3. Analyze Data
4. Visualize Results
5. Create interface with a tab called "Data" that displays the head of the dataset
6. Deploy to cloud

Two Main Deployment Options



Public deployment

You get a public URL
that anyone can visit



Private deployment

when your app needs
to stay private

What About Other Options?

👀 Other ways to deploy Streamlit apps

- Replit 
- Vercel 
- Netlify Drop 

👉 Recommendation

- ✓ Streamlit in Snowflake
- ✓ Streamlit Community Cloud



- Native Python support 
- Built-in security 

Deploying to Streamlit Community Cloud

The screenshot shows the 'App settings' page for a Streamlit app. On the left, there's a sidebar with 'App settings' (highlighted with a red box), 'General', 'Sharing', and 'Secrets'. The main area has two sections: 'Who can manage this app' (with a note about GitHub access) and 'Who can view this app'. The 'Who can view this app' section is also highlighted with a red box. It contains a dropdown menu set to 'This app is public and searchable', a note about specific viewing, and a list of email addresses ('person1@example.com', 'person2@example.com'). A 'Save changes' button is at the bottom.

 One private app per account



DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

Deploying your Prototype
Internally on Snowflake

Deploying your Prototype Internally on Snowflake

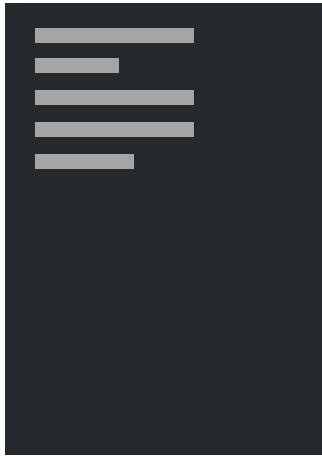


Download the sentiment analysis dashboard code from the GitHub repo:

```
M3
└ Lesson_01
    └ deploy
        └ streamlit_app.py
```

Updating Your App

Your app uses Snowflake's built-in connection methods



Authentication

Role-based access control (RBAC)

Automatic session management and connection pooling for SQL queries

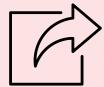
Role Based Access Control (RBAC)



Share your app with other users in your Snowflake account

- Access is controlled by existing Snowflake permissions and roles
- Users will need access to the underlying database and schema to view data
- The app remains secure within your Snowflake environment

Role Based Access Control (RBAC)



Share your app with other users in your Snowflake account



users
with
access



Database
└ Schema

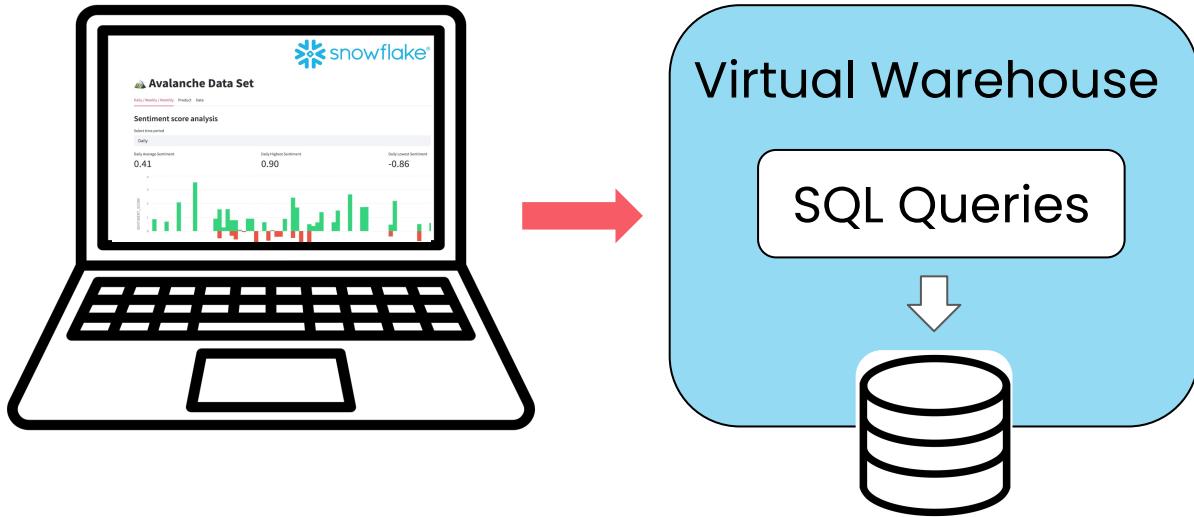
AVALANCHE_DB
AVALANCHE_SCHEMA



can view the
app

- Admin Role → modify app settings and permissions.
- Data Access → follows existing Snowflake security policies.

Monitoring Your App



Warehouse's size impacts:

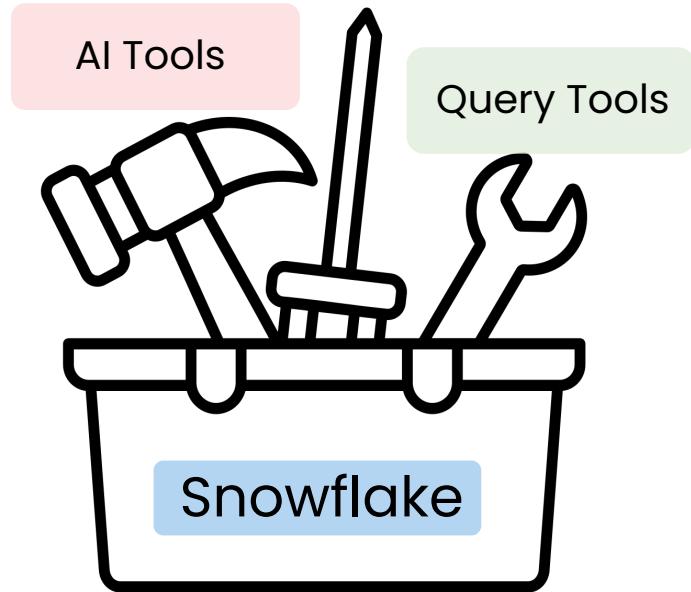
- Performance 
- Cost 

Key Metrics to Watch

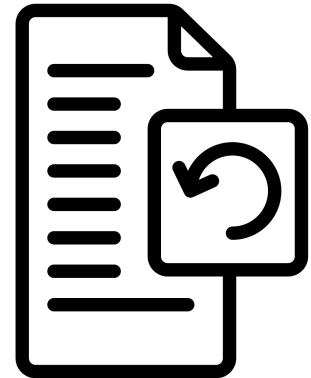


- How much of your **warehouse capacity** is being used?
- How long do individual **queries** take to complete?
- The **cost** of running your warehouse over time
- How many people are **using** your app simultaneously?

Checking Query History



Activity → Query History



Checking Warehouse Usage

❄️ > Admin > Warehouses



The Warehouses Tool shows both real-time and historical warehouse performance.

Resource Monitors → shows “cost control and usage alerts”

Practical Optimization Workflow

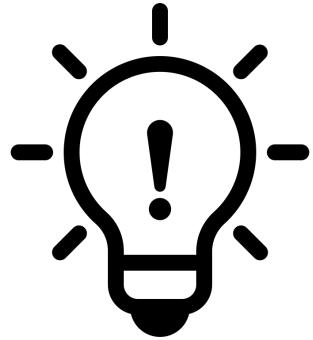


Weekly review process to monitor your app usage



- Check Query History
- Analyze Patterns
- Review Warehouse Usage
- Test Optimizations
- Monitor Results

Best Practices for Internal Deployment



- Provide clear navigation and instructions
- Include helpful tooltips and descriptions
- Ensure responsive design for different screen sizes
- Test with different user roles and permissions before rolling out



DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

Deploying to Streamlit Community Cloud

Join Streamlit Community Cloud

The screenshot shows the Streamlit Community Cloud landing page. At the top, there's a navigation bar with links for Playground, Gallery, Components, Cloud, Community, and Docs. A message "Introducing top navigation! See what's new in Release 1.46 🎉" is displayed above the navigation. To the right, there's a "Deploying? Try:" button with options for Free and Pro. The main heading "Community Cloud" is prominently displayed in large, bold, dark font. Below it, a sub-headline says "Share your apps with the whole world. Explore and fork community apps. Craft your profile. Totally free." Two buttons are present: "Join Community Cloud" (blue) and "Explore user apps" (white). A callout box titled "Deploy an app" contains instructions: "Apps are deployed directly from their GitHub repo. Enter the location of your app below." It features a "Repository" input field containing "streamlit-apps/repo" and a "Paste GitHub URL" link.

<https://streamlit.io/cloud>

Create Your Github Repo

1

Make a new Github Repo



Reminder: You need to choose a new distinct name for your repo

Upload Your Files

② Uploading files in your new repo

- ✓ In your new Github Repo, click on "**Upload an existing file**"
- ✓ When the upload is done, click on "**Commit changes**"

M3
└ Lesson_01
 └ deploy
 └ `streamlit_app.py`
 └ `requirements.txt`

Deploy to Streamlit Community Cloud

③ Deploy

Log in to
streamlit.io/cloud

Click “Create app”

Deploy a Public App from Github

 Repository, Branch: main

 Pointing out to the streamlit_app.py file

Choose a unique custom name

Click on “Deploy”!

Configure Snowflake Secrets

 Configure your credentials to access the Avalanche Snowflake database

-  Go to your deployed app
-  Click "**Manage App**" in the bottom-right corner
-  Click on the three dots → and then "**Settings**"
-  In the sidebar, click on "**Secrets**"

Configure Snowflake Secrets

```
[connections.snowflake]  
account = "xxxxxxxx-xxxxxxx"  
user = "your_username"  
password = "xxxxxxxxxx"  
role = "ACCOUNTADMIN"  
warehouse = "COMPUTE_WH"  
database = "AVALANCHE_DB"  
schema = "AVALANCHE_SCHEMA"
```

Configure Snowflake Secrets

https://<account_identifier>.snowflakecomputing.com

! Click on Save Changes!

If Things Go Wrong



- **App won't start?**
 - Check all uploaded files into GitHub
 - Make sure your `requirements.txt` file doesn't have any typos
- **Can't connect to Snowflake?**
 - Double-check
 - account identifier
 - Username and password



DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

Lab 1- Deploying Your Prototype

Part 1: Internal Deployment in Snowflake

① **Log
into
Snowsight**

Snowflake > Projects > Streamlit

Click the blue “**+ Streamlit App**” button

Name your app → “Avalanche Prototype”

Select: AVALANCHE_DB and
AVALANCHE_SCHEMA

Choose your compute warehouse

Click **Create**

Part 1: Internal Deployment in Snowflake

② Paste in Your Code

Editor: delete the default “Hello World” code

Copy from **streamlit_app.py** file and paste

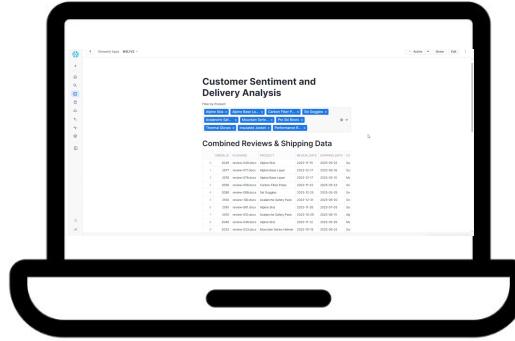
Make sure your code connects using:

`st.connection('snowflake')`

`get_active_session()`

Part 1: Internal Deployment in Snowflake

Click **Run** to launch your prototype



② **Run and Share Your App**

Click **Share** to get a link you can send to teammates with access

Monitor and Optimize

- ✓ Check the Query History
- ✓ Monitor Warehouse Usage to track cost, load, and idle time
- ✓ Set Resource Monitors to get alerts and manage spending



Pro tip:

Do a **weekly review**:

- Slow queries
- Redundant operations you can cache
- High-cost compute spikes



Bonus: Use Copilot



- Draft documentation
- Suggest code improvements
- Generate onboarding guides
- Fix performance issues

Part 2: Deploy to Streamlit Community Cloud

1 **Push to GitHub**

Make your app available to the world! 



Part 2: Deploy to Streamlit Community Cloud

1 **Push to GitHub**

Create a **new public** GitHub repo

Upload the following files:

-  streamlit_app.py
-  requirements.txt
-  streamlit/config.toml
-  any relevant data folders

Part 2: Deploy to Streamlit Community Cloud

1

**Push to
GitHub**

```
git init
git add .
git commit -m "Initial commit"
git branch -M main
git remote add origin
https://github.com/your-username/avalanche.git
git push -u origin main
```

Part 2: Deploy to Streamlit Community Cloud

② Log into Streamlit Community Cloud

Login to streamlit.io/cloud and login with GitHub

Click **Create app** and choose the repo to deploy

Point it to `streamlit_app.py`

Click on “Deploy”!

Part 2: Deploy to Streamlit Community Cloud

② **Log into
Streamlit
Community
Cloud**

App Management Menu > “Manage App”

Click on Settings > Click on Secrets

```
[connections.snowflake]  
account = "your-account-id"  
user = "your_username"  
password = "your_password"  
role = "ACCOUNTADMIN"  
warehouse = "COMPUTE_WH"  
database = "AVALANCHE_DB"  
schema = "PUBLIC"
```

Click on Save

Test and Monitor

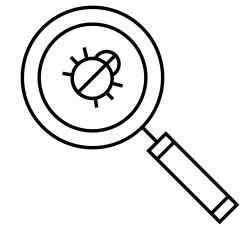
Visit the website address for your app:

<https://your-app-name.streamlit.app>

Settings menu >

click on **Manage App** >

Logs to check for any issues



Lesson 2

Getting Feedback



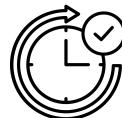
DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

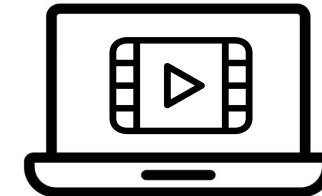
Iterate Quickly

Narrow Your Scope

With AI, you can build things faster than ever before!

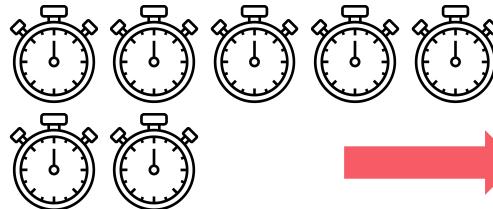


In short time



Build It Quick and Dirty, Then Fix It

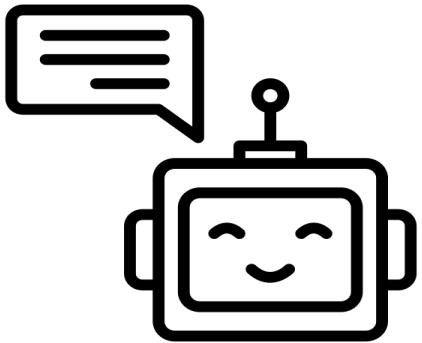
Traditional programming



GenAI programming



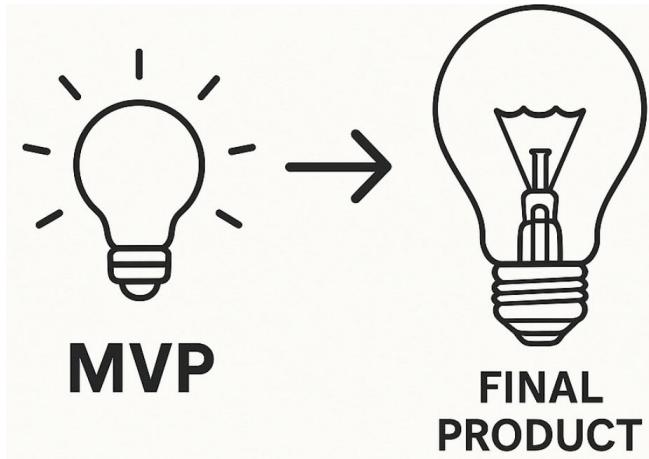
uncover problems by
programming



- Build a basic version of the chatbot
 - Answer the five common questions
- See where users get confused or frustrated
- Improve and repeat

Remember, It's Just a Prototype

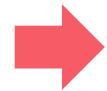
You're building something simple to test your ideas



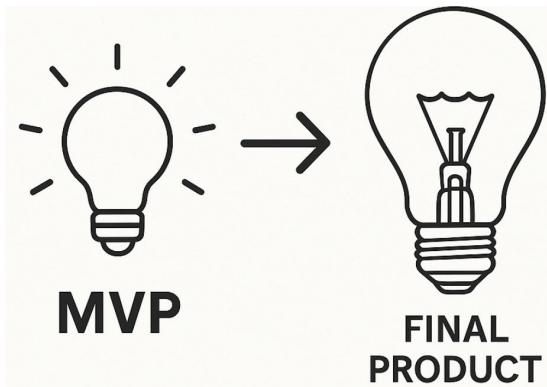
✓ The simplest
version

Know When to Stop

When is your prototype
"good enough"?



Done is better than perfect!

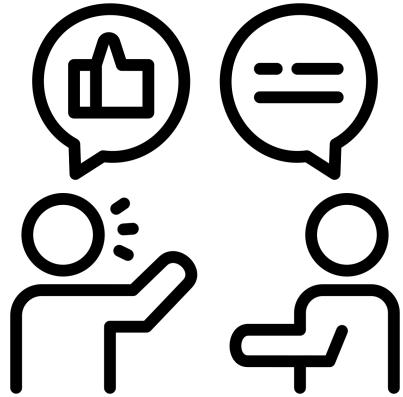


Your prototype just needs to:

- work well enough to test your main ideas
- get feedback

Deliver 80% of the core functionality, **efficiently**

Close The Loop With Users



- Tell people what you changed based on their feedback
- Send them updates
- Write release notes
- Show them in the app

Share What You Learned



Keep Improving



Development isn't a one-time thing



The people who succeed are the ones that keep listening and keep improving



Your competitors are getting better too, so you need to stay ahead



DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

Fast Feedback

Not All Feedback Is Created Equal

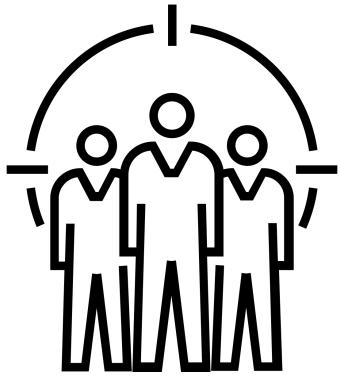


Early-stage AI development → qualitative feedback



Collect feedback efficiently and strategically

Start Small and Start Real



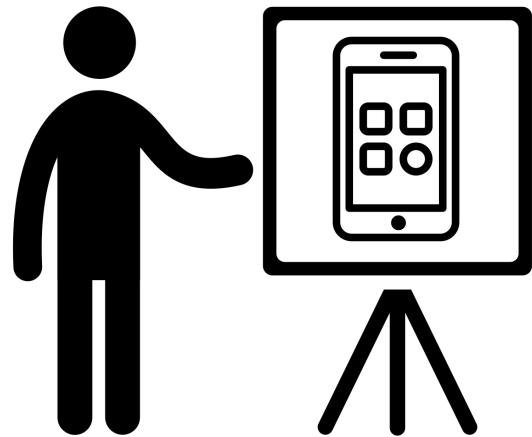
target user group



Find tester users:

- Social media
- Professional networks
- Existing customer base

The 5 Second Rule



Ask them:

- What do they remember?
- What do they think the app is supposed to do?

Ask at the Right Moment

Before Use

Low feedback
quality

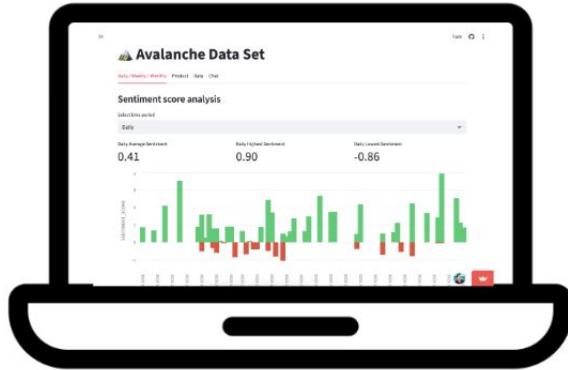
During Use

Distracted,
partial feedback

 **After Use**

Best time — fresh,
clear experience

The One True Path



Important →

Perform a sentiment analysis
on an input dataset



Ask your testers to:

- Select a date range and set of products, then review the results

Ask Better Questions

Broad & vague

- 🚫 "What do you think?"
- 🚫 "Do you like it?"

Specific & actionable

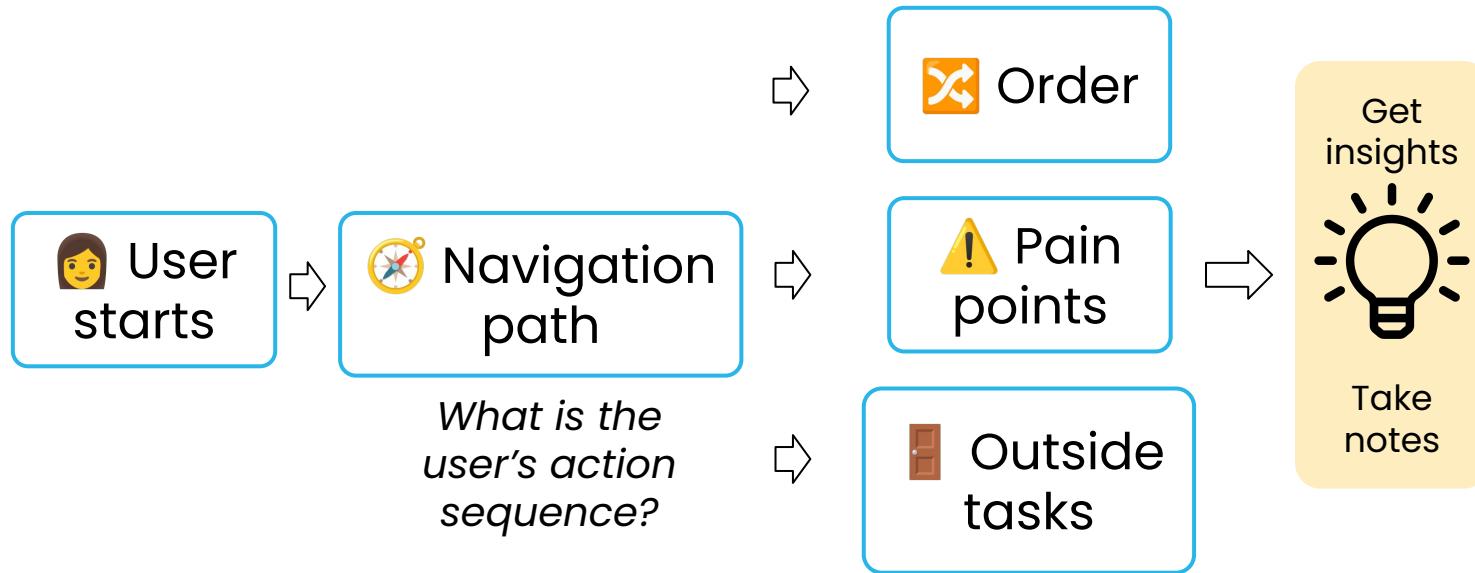
- ✓ "Was this button easy to find?"
- ✓ "How did this feature help you solve your problem?"



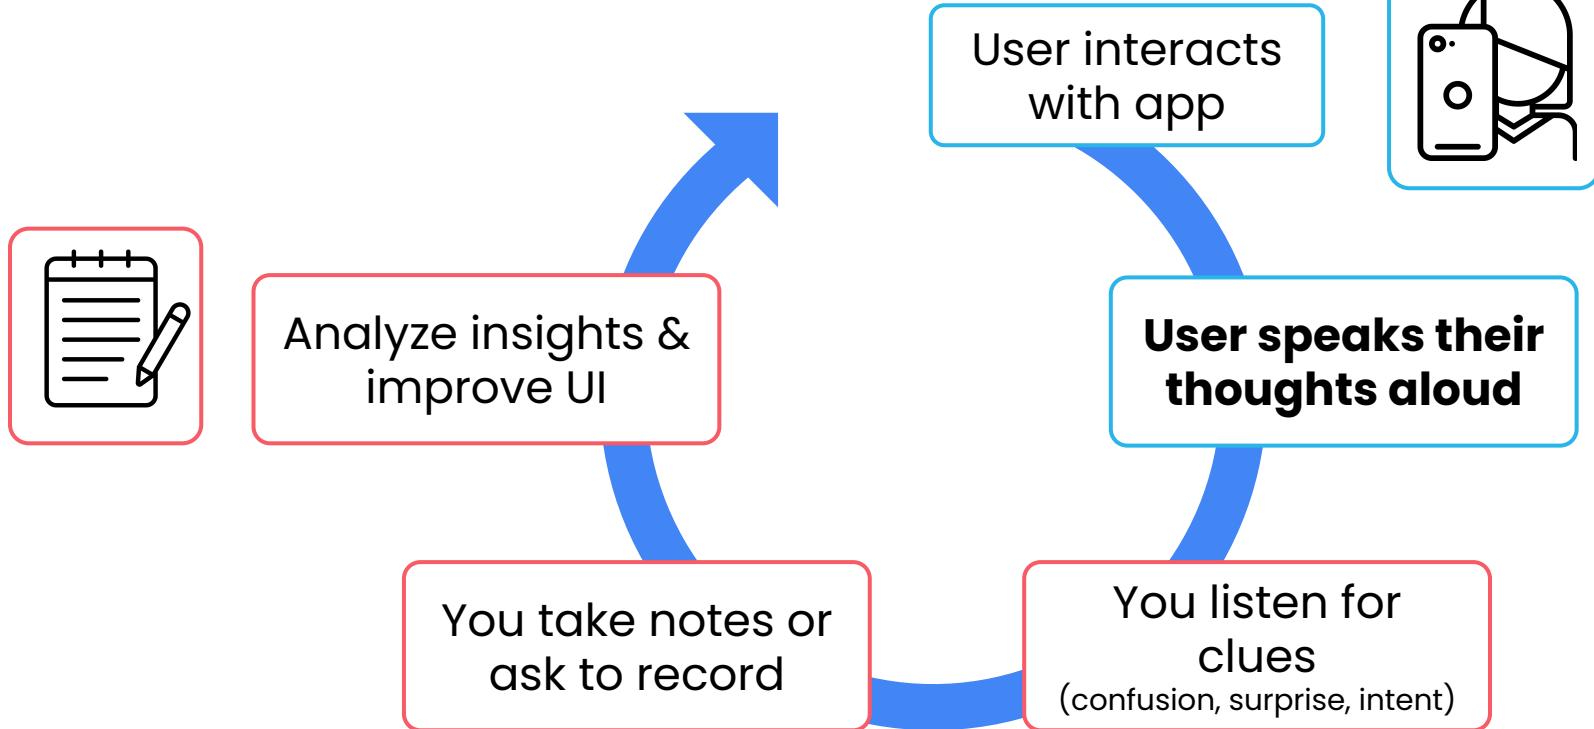
Ask open-ended questions → 🧠 Insights

Ask open-ended questions → 📊 Trends

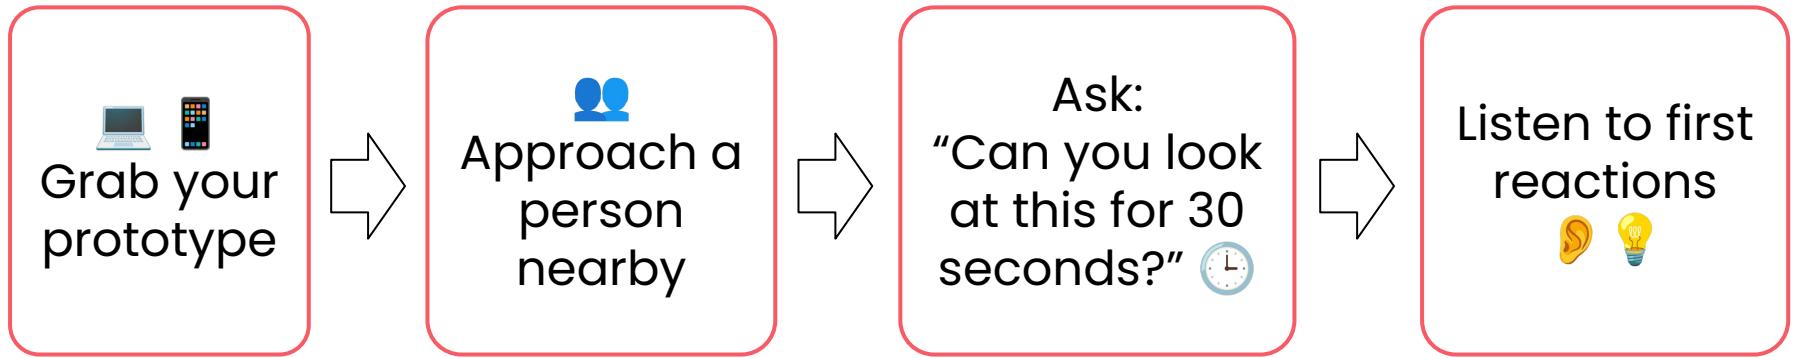
Channel Your Inner Scientist



Think Out Loud



Hallway Feedback



Wrap Up

- Gather feedback early and often 🎁
- Prioritize feedback using the impact vs. effort matrix ✋
- Focus on high-impact, low-effort changes first 🔎

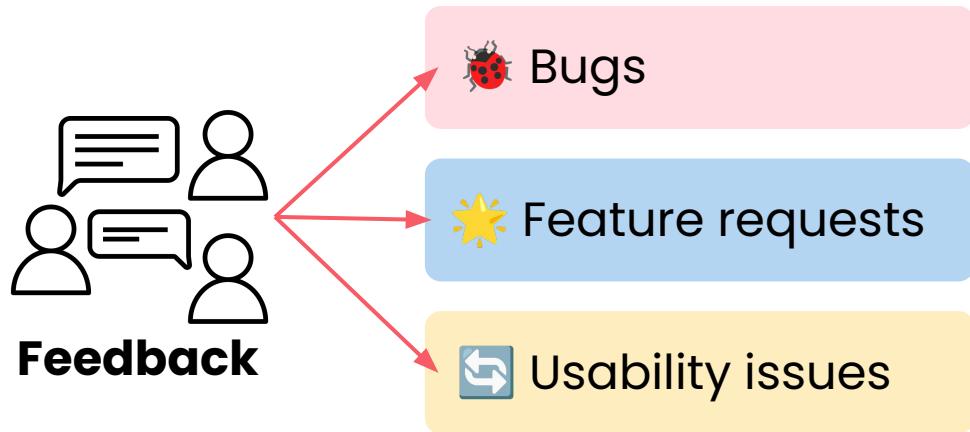


DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

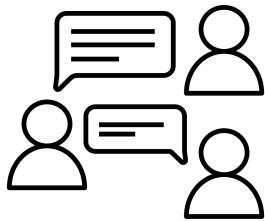
Acting on Feedback

Organize Feedback Into Buckets



- ✓ WHAT people are asking for
- ✓ WHY they are asking for it

Think Like a Detective



Feedback

"your app is slow"

"this feature doesn't work"



Think like a
detective



Break the
problem into
slices

- Is it slow for new users or returning users?
- Is it affecting mobile users more than desktop users?

Measure Your Changes

Your measurement approach should be just as fast as your building

 Quick validation methods



User observation



Micro-survey



Basic analytics check

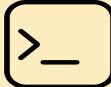


A/B test (simple)

Use AI To Speed Everything Up



Let AI help you code



Get good at prompting

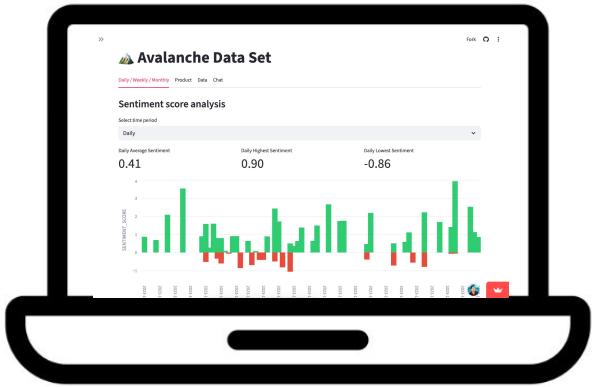


Try agentic workflows



Track feedback systematically

Acting on Feedback



Some comments:

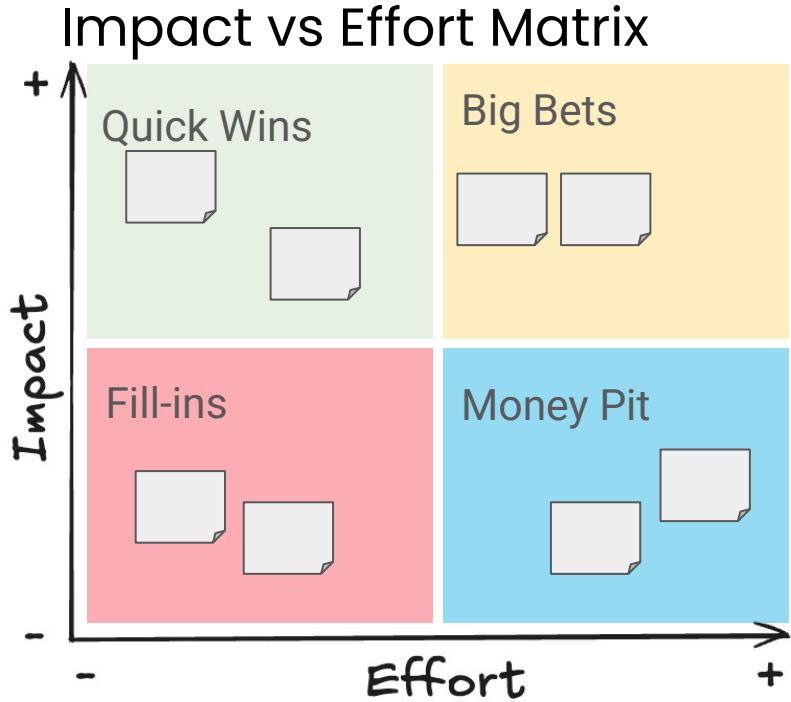
I asked it questions about the dataset, and it gave me incorrect answers

It didn't really seem to know much about our company's research methodology.

It would be great if we could choose other datasets to analyze.

Feedback Driven Prototyping

💡 Remember the
“Must Have”
“Nice to Have”
decision framework
→ MVP Building Plan

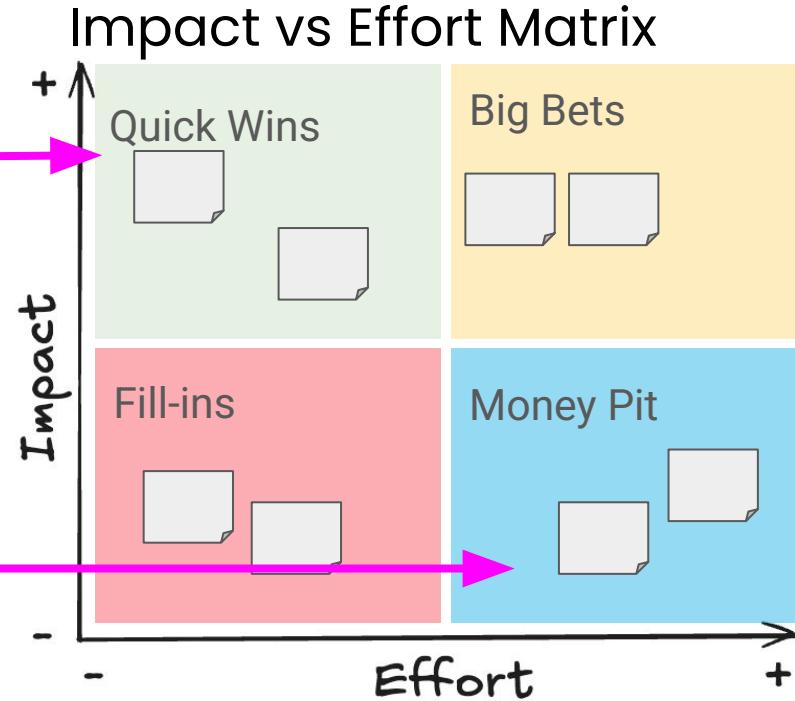


Feedback Driven Prototyping

I asked it questions about the dataset, and it gave me incorrect answers

It didn't really seem to know much about our company's research methodology.

It would be great if we could choose other datasets to analyze.





Fast Prototyping of GenAI Apps with Streamlit

Iterate, Improve, Repeat
– Fast Feedback for
Your Avalanche App

By The End of This Lesson

- ① Test your app from a user's perspective
- ② Collect actionable feedback quickly
- ③ Prioritize improvements based on impact and effort
- ④ Use AI tools to accelerate your iteration process

The Scenario



Open your deployed Avalanche app:

- Snowflake's Streamlit interface
- Streamlit Cloud (if you deployed there)

What do customers think about our winter product line?

Step 1: Critical User Journey Test



Evaluate:

- Can you quickly locate the sentiment analysis feature?
- Are the results presented in an easy-to-understand format?
- Would someone unfamiliar with your dataset know what to click next?

Step 2: Run a 5-Second Test



?

Ask:

- What do you think this app does?
- What stood out to you most?
- Could you figure out what to do next?

Step 3: Test the “One True Path”



User's journey:

- Load the dashboard
- Run sentiment analysis
- Interpret the results

? Ask yourself:

- Is anything confusing or awkward in this flow?
- Are instructions clear and actionable?
- Does the layout support this main task effectively?

Step 4: Organize Your Feedback



Bugs



Usability Issues



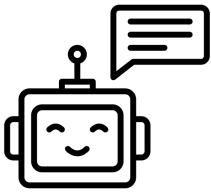
Feature Requests



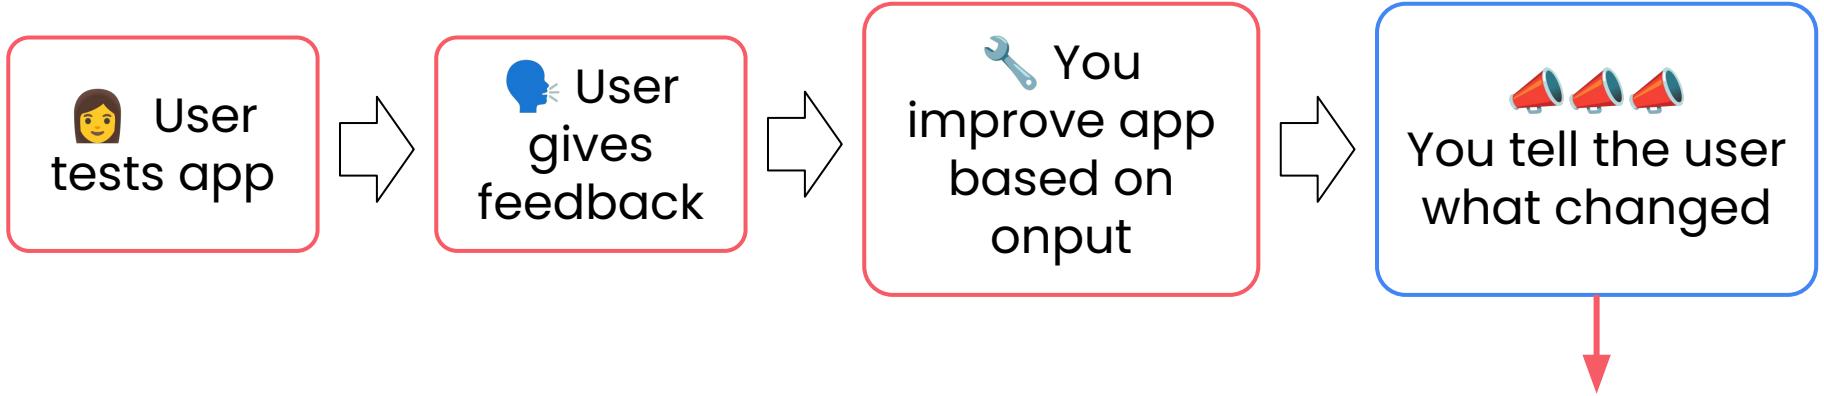
Ask:

- What's high-impact but low-effort to fix?
- What can I improve right now in under 30 minutes?

Step 5: Use AI to Accelerate Improvements

Tools	Content Issues	UX Problems
 GitHub Copilot, ChatGPT, Claude Ask AI?	rephrase confusing labels or tooltips	GenAI Assistant tab in Snowflake → test new ways of answering user questions
	rewrite prompt templates for better clarity	suggest alternative chart types or data visualizations
	cleaning up code blocks	

Step 6: Close the Loop



*"Hey, I made that chart easier to read **based on your feedback**. Thanks again for the tip!"*

Lesson 3

Improving Your Results



DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

Improving Prompts

Why Prompt Engineering Matters for Fast Prototyping

In fast prototyping, time is everything

- ✓ Build a working prototype quickly
- ✓ Test it with real users
- ✓ Gather feedback
- ✓ Implement changes fast
- ✓ Repeat the cycle

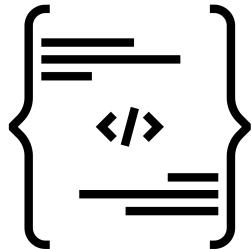
✗ Poor prompts slow you down

- ✓ A clear, specific prompt gets you 95%+ of the way

The Feedback Integration Problem



- "the chart is confusing"
- "I can't find the export button"



Try prompting your AI like this:

You are a UX-focused Python developer. The user said "the chart is confusing."

Rewrite this chart code to include:

- Clear axis labels
- A legend that explains what each line represents
- Hover tooltips showing exact values
- A title that explains what the chart shows

The Four Building Blocks of Effective Prompts



Role



Task

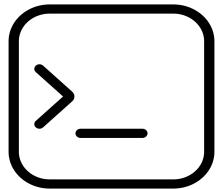


Output



Context

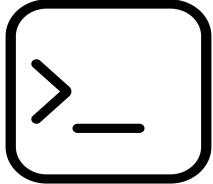
Instructions: Tell the AI Exactly What to Do



Make an app



hours modifying
generic code



"Build a Streamlit app that analyzes PDF research papers and answers questions about them"

Context: Provide Essential Background

🚫 Without context: AI might make wrong assumptions

✓ Add context for speed & precision

Role setting



"You are a Python developer experienced in Streamlit apps"

Constraints and requirements



Specify...

- Python versions
- deployment platforms
- technical limitations

Input Data: Supply the Raw Materials

🚫 Vague input → Generic output

✓ Real Input → Real solutions

Examples of concrete input:

- Specific code snippets to modify or extend
- Example data formats (e.g., CSV, JSON)
- Sample files or datasets

Output Indicators: Specify the Response Format

! Problem: Unclear output = Wasted time

✓ Solution: Be explicit about the format you want

… Say things like:

- "Return only working Python code"
- "Format the output as JSON"
- "Give me a single Python file"

🔄 When iterating fast:

- "Return only the modified function"
- "Show just the CSS changes"

Prompt Engineering Strategies for Fast Prototyping

① Start simple, then iterate

- Get something working quickly
- Improve it based on user feedback
- Don't try to build everything at once
- Create a basic version first, then refine it step by step

Benefits:

- ✓ Better control
- ✓ Reliable improvements
- ✓ Fast feedback cycles

Prompt Engineering Strategies for Fast Prototyping

② Break complex tasks into smaller prompts

- Basic app structure
- Add a specific feature
- Handle errors
- Optimize performance

Example:

 User: "App is slow & charts are confusing"

 Prompt 1:

"Optimize this data loading code for better performance"

 Prompt 2:

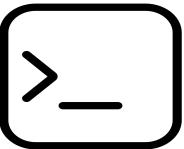
"Improve these chart visualizations for better clarity"

Use Guardrails for Consistency

Guardrails keep your code consistent and maintainable

Guardrail	Why It Helps
Tell the AI what to avoid	Prevents unwanted complexity or breaking changes
Ask for reasoning	Understand why changes are made — stay in control
Use consistent terminology	Keeps codebase coherent across iterations
Maintain your coding standards	Code stays clean and readable, even under pressure

A Real World Example



You are a Python developer experienced in Streamlit apps.

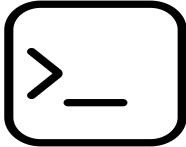
Build a PDF research paper analyzer with these features:

- Upload PDF functionality
- Extract and display title and abstract
- Summarize main contributions (under 300 words)
- Interactive Q&A using GPT-4

Assume PyMuPDF and Streamlit are installed.

Return only complete working code in a single Python file.

The Role Setting



You are a Python developer experienced in Streamlit apps.

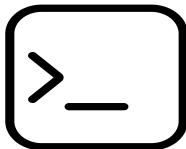
Build a PDF research paper analyzer with these features:

- Upload PDF functionality
- Extract and display title and abstract
- Summarize main contributions (under 300 words)
- Interactive Q&A using GPT-4

Assume PyMuPDF and Streamlit are installed.

Return only complete working code in a single Python file.

Clear Task + Features



You are a Python developer experienced in Streamlit apps.

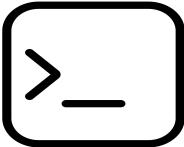
Build a PDF research paper analyzer with these features:

- Upload PDF functionality
- Extract and display title and abstract
- Summarize main contributions (under 300 words)
- Interactive Q&A using GPT-4

Assume PyMuPDF and Streamlit are installed.

Return only complete working code in a single Python file.

The Technical Assumptions



You are a Python developer experienced in Streamlit apps.

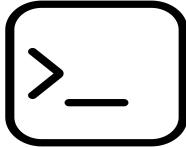
Build a PDF research paper analyzer with these features:

- Upload PDF functionality
- Extract and display title and abstract
- Summarize main contributions (under 300 words)
- Interactive Q&A using GPT-4

Assume PyMuPDF and Streamlit are installed.

Return only complete working code in a single Python file.

The Output Specification



You are a Python developer experienced in Streamlit apps.

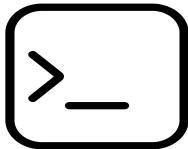
Build a PDF research paper analyzer with these features:

- Upload PDF functionality
- Extract and display title and abstract
- Summarize main contributions (under 300 words)
- Interactive Q&A using GPT-4

Assume PyMuPDF and Streamlit are installed.

Return only complete working code in a single Python file.

Refining Apps with GenAI Feedback



You are a UX-focused Python developer. Based on user feedback that the dashboard is "too cluttered" and they "can't find the export button," modify this Streamlit code to:

- Move the export button to a prominent location in the sidebar
- Reduce visual clutter by using expandable sections for secondary features
- Add clear section headers to organize the content
- Ensure the most important metrics are visible at the top

Return only the modified code sections.

Review and Enhancement Process

Systematic approach to keep improving your prototype:

- Rapid Code Review 
- Performance Optimization 
- Feature Expansion 
- UX Enhancements 

Prompt → Prototype → Feedback → Improvement → Repeat



DeepLearning.AI

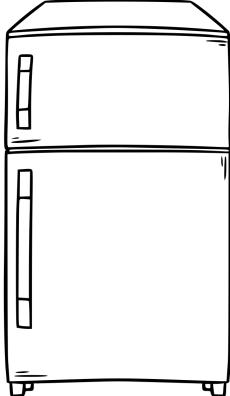
Fast Prototyping of GenAI Apps with Streamlit

Upgrading Your Prototype with Data Augmentation

What is Prompt Augmentation?

Add specific information about your dataset right into the prompt

Give your prototype → a cheat sheet



“Guess what’s in my fridge.”



“Here’s what’s in the fridge –
now tell me what I can cook”

Gen AI Chatbot and Avalanche

: "What products are in the Avalanche dataset?"

: "I don't have access to specific product information."

*Currently, the chatbot knows nothing about the Avalanche dataset 😞

 Solution



Feed the Data into the Prompt

Getting Set Up

Starting Point:



→ Located in the GitHub repo: M3/Lesson_03/



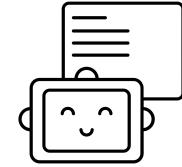
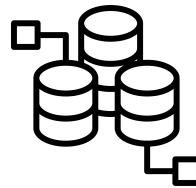
DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

Using RAG to Improve Model Performance

What is RAG?

RAG: Retrieval-Augmented Generation



User query



Retrieval system



Generation



Final answer

- Search docs or DB
- Find top relevant info

- Input info + query
- LLM writes a response

Why Does This Matter?

🚫 LLMs don't know everything

- Don't know your latest product data
- Can't access internal documents or company policies
- Aren't reliable in niche domains (e.g., medicine, law)

✓ RAG = smarter responses without retraining

- Pulls info at runtime
- Ideal for accuracy, transparency, and trust
- Great for apps in:
 - medical research, 
 - contract law, 
 - enterprise data 

When Use RAG Instead of Just Prompting?

⚠ Challenge

Your knowledge base is too large to fit in a prompt → Retrieve only the most relevant info

Your content changes frequently

→ Pulls real-time data at the moment

You want source attribution

→ Keeps links or citations

Working in specialized field

→ Boosts accuracy

You want personalized responses

→ Fetches user-specific data on the fly

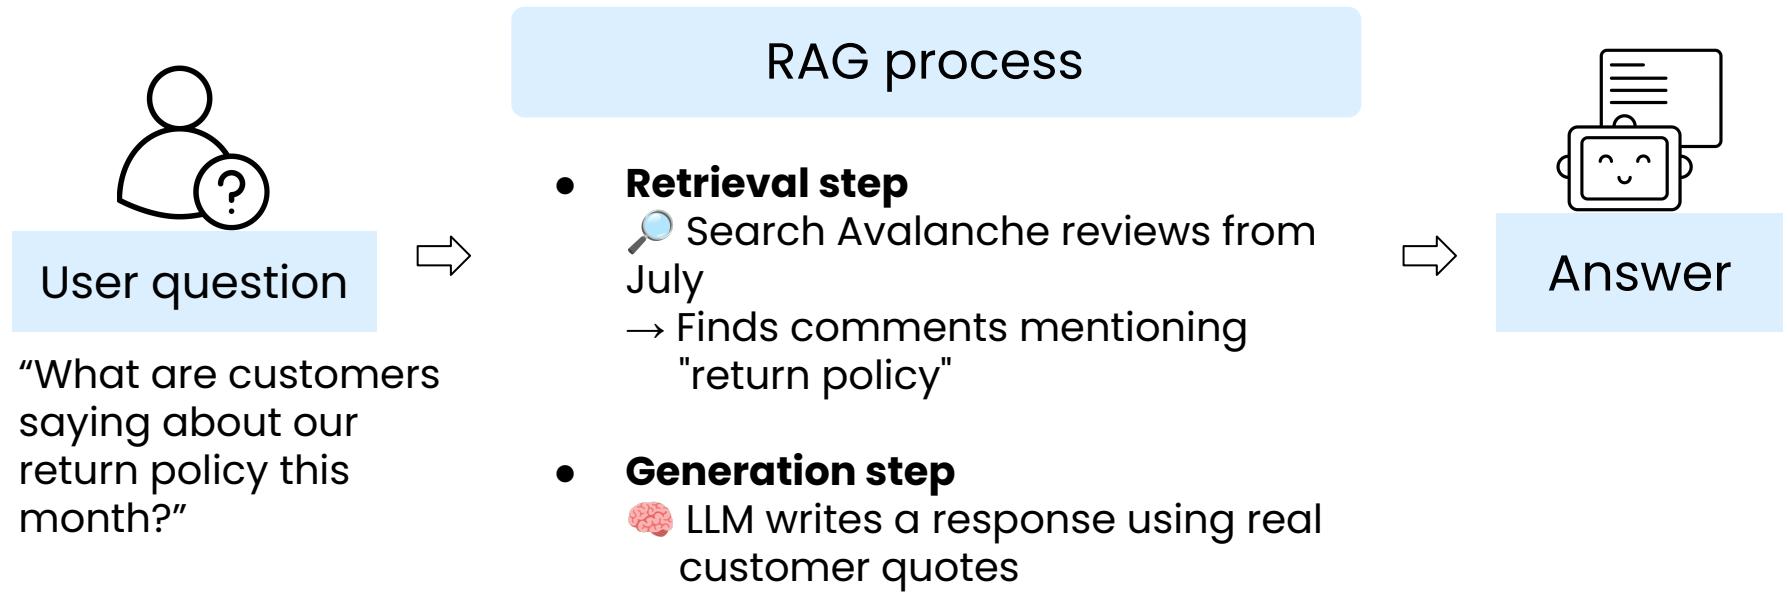
You're optimizing token usage

→ Smaller prompts = lower cost

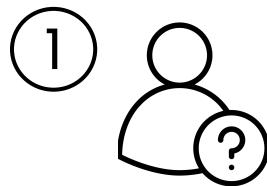


RAG Solution

Let's Bring it Back to the Avalanche prototype



Working Behind The Scenes



User question

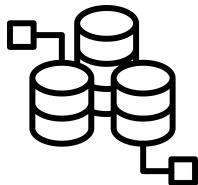
- 1
- 2 Break data into chunks
- 3 Retrieve relevant chunks
- 4 Feed chunks + question to LLM
- 5 Generate grounded answer

Snowflake Cortex

<https://docs.snowflake.com/en/user-guide/snowflake-cortex/cortex-search/cortex-search-overview>

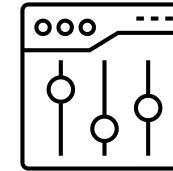
What About Fine-tuning? How is RAG Different?

RAG



New context at query time

Fine-Tuning



Behavior into the model itself

- Need a specific writing style or tone
- Doing the same task over and over
- Need consistent logic or reasoning
- Don't want to manage external documents or storage

Most Production Systems Use a Hybrid Approach

Prompt engineering

⇒ guide the model step-by-step

RAG

⇒ inject relevant, factual context

Fine-tuning

⇒ lock in style, behavior, or reasoning for repetitive tasks



DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

Setting up a RAG Pipeline Using Cortex Search

Setting Up a RAG Pipeline Using Cortex Search

- Cortex Search is Snowflake's managed search service for AI applications
 -  Automatic indexing & embeddings
 -  Simple APIs for RAG applications
 -  No infrastructure management required

Setting Up a RAG Pipeline Using Cortex Search

Follow along by:

- Opening a new Snowflake notebook
- Or finding the code in the course GitHub repo:
 - M3/Lesson_03/Lab2



DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

Lab 2- Integrating RAG into Your Chatbot

You Did It!

- ✓ From basic Streamlit to advanced RAG integration.
- ✓ You can quickly prototype any data app idea, add GenAI capabilities, and turn concepts into working solutions.



DeepLearning.AI

Fast Prototyping of GenAI Apps with Streamlit

What Comes Next?

What You've Learned

- From ingestion → to analysis → to UX
- Using real-world data
- Streamlit to build polished front ends, fast
- Snowflake for powerful data handling
- Cortex to supercharge it all with GenAI

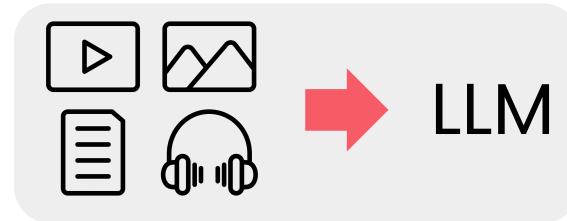


Working prototype



What Comes Next?

① **Build a
multi-modal app**



"You record a
screencast of
your app in use..."



Feed it to a
model...



Ask for

- UX feedback
- Documentation
- Next-step features – all automatically

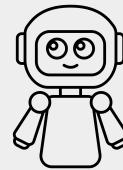
What Comes Next?

②

Prototype with AI agents



Create custom agents that behave like junior developers



agentic AI

Give:

- persona
- some goals
- dev style ...

MVP-first approach

What Comes Next?

Turn your ③ prototype into a product



For work,
a side project,
or just for fun

idea → working GenAI app

prototype → product

- Add authentication and permissions
- Integrate with external APIs
- Turn insights into features
- Deploy and monetize

The Mindset That Matters

- AI will keep evolving 🔥
- The tools will change 🚧
- The interfaces will change 💻
- The models will change 📊



Mindset!



Ideate. Build. Evaluate. Iterate. Improve. Repeat.

Concluding Remarks

Got the skills to build smarter apps – faster:

- Keep prototyping
- Keep experimenting
- Keep letting GenAI amplify what you bring to the table



Congratulations!